

# Lecture 14 Notes

2024-03-05

## Opening the data

```
require(tidyverse)
```

```
## Loading required package: tidyverse
```

```
## Warning: package 'tidyverse' was built under R version 4.3.2
```

```
## — Attaching core tidyverse packages — tidyverse 2.0.0 —  
## ✓ dplyr      1.1.2      ✓ readr      2.1.4  
## ✓ forcats    1.0.0      ✓ stringr    1.5.0  
## ✓ ggplot2    3.4.4      ✓ tibble     3.2.1  
## ✓ lubridate  1.9.2      ✓ tidyr      1.3.0  
## ✓ purrr      1.0.1
```

```
## — Conflicts — tidyverse_conflicts() —  
## ✗ dplyr::filter() masks stats::filter()  
## ✗ dplyr::lag()     masks stats::lag()  
## ⓘ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to  
o become errors
```

```
mv <- read_rds('https://github.com/jbisbee1/DS1000_S2024/raw/main/data/mv.Rds')  
  
mv_analysis <- mv %>%  
  drop_na(budget, gross) %>%  
  mutate(budget_log = log(budget),  
         gross_log = log(gross))
```

## Regression

```
model_gross_budget <- lm(formula = gross_log ~ budget_log,  
                          data = mv_analysis)
```

## RMSE

```

mv_analysis <- mv_analysis %>%
  mutate(preds = predict(model_gross_budget)) %>%
  mutate(errors = gross_log - preds)

rmse <- mv_analysis %>%
  mutate(se = errors^2) %>%
  summarise(mse = mean(se)) %>%
  mutate(rmse = sqrt(mse))

```

## Cross Validation

```

set.seed(123)
cvRes <- NULL
for(i in 1:100) {
  # Generating a random list of rows
  inds <- sample(x = 1:3179,
    size = round(3179/2),
    replace = FALSE)

  train <- mv_analysis %>%
    slice(inds)

  test <- mv_analysis %>%
    slice(-inds)

  # Training on the training data
  model_tmp <- lm(gross_log ~ budget_log, data = train)

  # Predicting on the test data
  test <- test %>%
    mutate(preds = predict(model_tmp,
      newdata = test)) %>%
    mutate(errors = gross_log - preds)

  rmse <- test %>%
    mutate(se = errors^2) %>%
    summarise(mse = mean(se)) %>%
    mutate(rmse = sqrt(mse))

  cvRes <- cvRes %>%
    bind_rows(rmse)
}

cvRes %>%
  summarise(rmseOverall = mean(rmse))

```

```
## # A tibble: 1 × 1
##   rmseOverall
##       <dbl>
## 1         1.28
```


## Alternative model

```
mv_analysis %>%
  select(gross_log, score)
```

```
## # A tibble: 3,179 × 2
##   gross_log score
##       <dbl> <dbl>
## 1      18.1   7.9
## 2      17.8   7.6
## 3      20.4   8.5
## 4      16.3   8.3
## 5      17.9   8.4
## 6      20.3   7.8
## 7      19.9   6.2
## 8      20.1   6.4
## 9      19.0   5.7
## 10     19.9   7.4
## # i 3,169 more rows
```

```
# Univariate viz
mv_analysis %>%
  ggplot(aes(x = score)) +
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
# Missingness
summary(mv_analysis %>%
  select(gross_log,score))
```

```
##   gross_log      score
##  Min.   : 8.287   Min.   :1.900
##  1st Qu.:17.068   1st Qu.:5.900
##  Median :18.189   Median :6.500
##  Mean   :17.922   Mean   :6.417
##  3rd Qu.:19.104   3rd Qu.:7.100
##  Max.   :21.991   Max.   :9.000
```

```
# Multivariate visualization
mv_analysis %>%
  ggplot(aes(x = score,y = gross_log)) +
  geom_point() +
  geom_smooth(method = 'lm')
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



Cross Validation: Model 2

```

set.seed(123)
cvRes <- NULL
for(i in 1:100) {
  # Generating a random list of rows
  inds <- sample(x = 1:nrow(mv_analysis),
                size = round(nrow(mv_analysis)/2),
                replace = FALSE)

  train <- mv_analysis %>%
    slice(inds)

  test <- mv_analysis %>%
    slice(~inds)

  # Training on the training data
  model_tmp <- lm(gross_log ~ score, data = train)

  # Predicting on the test data
  test <- test %>%
    mutate(preds = predict(model_tmp,
                           newdata = test)) %>%
    mutate(errors = gross_log - preds)

  rmse <- test %>%
    mutate(se = errors^2) %>%
    summarise(mse = mean(se)) %>%
    mutate(rmse = sqrt(mse))

  cvRes <- cvRes %>%
    bind_rows(rmse)
}

cvRes %>%
  summarise(rmseOverall = mean(rmse))

```

```

## # A tibble: 1 × 1
##   rmseOverall
##         <dbl>
## 1         1.75

```

## Multiple X Variables

```

#  $Y = \alpha + \beta_1 X_1 + \beta_2 X_2 + \dots + e$ 
model_budg_score <- lm(gross_log ~ budget_log + score,
                      mv_analysis)

summary(model_budg_score)

```

```
##
## Call:
## lm(formula = gross_log ~ budget_log + score, data = mv_analysis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.3763 -0.5708  0.1501  0.7322  8.6189
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.70352    0.33839  -2.079   0.0377 *
## budget_log   0.96710    0.01742  55.527  <2e-16 ***
## score        0.29740    0.02316  12.843  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.249 on 3176 degrees of freedom
## Multiple R-squared:  0.5041, Adjusted R-squared:  0.5038
## F-statistic: 1614 on 2 and 3176 DF, p-value: < 2.2e-16
```

```
(exp(0.29740)-1)*100
```

```
## [1] 34.63537
```

## Cross Validation: Model 3

```

set.seed(123)
cvRes <- NULL
for(i in 1:100) {
  # Generating a random list of rows
  inds <- sample(x = 1:nrow(mv_analysis),
                size = round(nrow(mv_analysis)/2),
                replace = FALSE)

  train <- mv_analysis %>%
    slice(inds)

  test <- mv_analysis %>%
    slice(~inds)

  # Training on the training data
  model_tmp <- lm(gross_log ~ budget_log + score, data = train)

  # Predicting on the test data
  test <- test %>%
    mutate(preds = predict(model_tmp,
                          newdata = test)) %>%
    mutate(errors = gross_log - preds)

  rmse <- test %>%
    mutate(se = errors^2) %>%
    summarise(mse = mean(se)) %>%
    mutate(rmse = sqrt(mse))

  cvRes <- cvRes %>%
    bind_rows(rmse)
}

cvRes %>%
  summarise(rmseOverall = mean(rmse))

```

```

## # A tibble: 1 × 1
##   rmseOverall
##         <dbl>
## 1         1.25

```

## Movie Ratings

```

model_gross_rating <- lm(gross_log ~ rating,
                        mv_analysis)

summary(model_gross_rating)

```



```
##
## Call:
## lm(formula = gross_log ~ rating, data = mv_analysis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.6749 -0.8191  0.1630  1.1082  5.2339
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    19.1818     0.2182  87.914 < 2e-16 ***
## ratingNC-17     -2.4483     0.6900  -3.548 0.000393 ***
## ratingNot Rated -4.4322     0.3510 -12.626 < 2e-16 ***
## ratingPG        -0.3905     0.2314  -1.688 0.091517 .
## ratingPG-13     -0.7633     0.2229  -3.425 0.000623 ***
## ratingR         -1.9123     0.2224  -8.599 < 2e-16 ***
## ratingTV-MA     -3.2064     1.1546  -2.777 0.005515 **
## ratingUnrated   -4.6564     0.6441  -7.229 6.05e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.603 on 3166 degrees of freedom
## (5 observations deleted due to missingness)
## Multiple R-squared:  0.1771, Adjusted R-squared:  0.1753
## F-statistic: 97.33 on 7 and 3166 DF, p-value: < 2.2e-16
```

## Movie Genre

```
model_gross_genre <- lm(gross_log ~ genre,
                        mv_analysis)

summary(model_gross_genre)
```

```
##
## Call:
## lm(formula = gross_log ~ genre, data = mv_analysis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.8819 -0.7452  0.2469  1.0992  3.6548
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   18.58333    0.05560  334.235 < 2e-16 ***
## genreAdventure -0.18471    0.13449   -1.373  0.169712
## genreAnimation  0.67982    0.12467    5.453 5.33e-08 ***
## genreBiography -1.14278    0.12421   -9.200 < 2e-16 ***
## genreComedy    -1.00044    0.08101  -12.350 < 2e-16 ***
## genreCrime     -1.39688    0.12729  -10.974 < 2e-16 ***
## genreDrama     -1.41465    0.09161  -15.442 < 2e-16 ***
## genreFamily     0.37140    1.16627    0.318  0.750166
## genreFantasy   -1.69110    0.46030   -3.674  0.000243 ***
## genreHorror    -0.91636    0.14597   -6.278 3.90e-10 ***
## genreMystery    0.03469    0.62516    0.055  0.955749
## genreRomance   -2.68759    0.82561   -3.255  0.001145 **
## genreSci-Fi    -0.16302    1.16627   -0.140  0.888846
## genreThriller  -0.58577    0.82561   -0.709  0.478068
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.647 on 3165 degrees of freedom
## Multiple R-squared:  0.1407, Adjusted R-squared:  0.1372
## F-statistic: 39.87 on 13 and 3165 DF, p-value: < 2.2e-16
```

```
mv_analysis %>%
  count(genre)
```

```
## # A tibble: 14 × 2
##   genre      n
##   <chr>    <int>
## 1 Action    878
## 2 Adventure 181
## 3 Animation 218
## 4 Biography 220
## 5 Comedy    782
## 6 Crime      207
## 7 Drama      512
## 8 Family       2
## 9 Fantasy     13
## 10 Horror    149
## 11 Mystery     7
## 12 Romance     4
## 13 Sci-Fi      2
## 14 Thriller    4
```

## Bechdel Score

```
model_gross_BS <- lm(gross_log ~ bechdel_score,
                     mv_analysis)
```

```
summary(model_gross_BS)
```

```
##
## Call:
## lm(formula = gross_log ~ bechdel_score, data = mv_analysis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.9000 -0.7755  0.2216  1.0792  3.8042
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  18.36809    0.08565  214.443  <2e-16 ***
## bechdel_score -0.06042    0.03545   -1.704   0.0885 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.665 on 2056 degrees of freedom
## (1121 observations deleted due to missingness)
## Multiple R-squared:  0.001411,    Adjusted R-squared:  0.0009253
## F-statistic: 2.905 on 1 and 2056 DF,  p-value: 0.08845
```

```
mv_analysis %>%  
  drop_na(bechdel_score) %>%  
  ggplot(aes(x = factor(bechdel_score),  
             y = gross_log)) +  
  geom_violin()
```

