

Classification

Part 3

Prof. Bisbee

Vanderbilt University


Slides Updated: 2024-03-16

Agenda

1. Recap of regression and classification
2. Introducing (some) machine learning algorithms

What is regression?


- Conditional means for continuous data




Regression

- Calculating a **line** that minimizes mistakes *for every observation*
 - NB: could be a curvey line! For now, just assume straight
- Recall from geometry how to graph a straight line
- $Y = a + bX$
 - a : the "intercept" (where the line intercepts the y-axis)
 - b : the "slope" (how much Y changes for each increase in X)
- (Data scientists use α and β instead of a and b b/c nerds)
- Regression analysis simply chooses the best line
 - "Best"?
 - The line that minimizes the mistakes (the **line of best fit**)

Visual Intuition



Visual Intuition




Two Camps Revisited

- Regression is great for **theory testing**
 - Results tell us something **meaningful** about our theory
- But if all we care about is **prediction...**?
 - Want to test every possible predictor (and combinations)
 - Don't care about **relationships**
 - Just care about **accuracy**
- Algorithms can save us time!
 - Random Forests
 - LASSO

Random Forests

- Identify the best "partition" (split) that divides the data



- In R: `ranger`
 - `formula = Y ~ .`

Random Forests

```
require(tidyverse)
require(scales)
require(tidymodels)
fn <-
read_rds('C:/Users/bisbeejh/Dropbox/2024_spring/DS1000_S2024/data/fn_cl
```

Research Question

- What predicts whether you win at Fortnite?

```
form.perf <- 'won ~ hits + assists + accuracy + head_shots +  
damage_to_players'
```

```
form.games <- 'won ~ eliminations + revives + distance_traveled +  
materials_gathered'
```

```
form.context <- 'won ~ mental_state + startTime + gameIdSession'
```

```
form.full <- 'won ~ hits + assists + accuracy + head_shots +  
damage_to_players + eliminations + revives + distance_traveled +  
materials_gathered + mental_state + startTime + gameIdSession'
```

Comparing models

```
m.perf <- lm(as.formula(form.perf),fn)
summary(m.perf)
```

```
##
## Call:
## lm(formula = as.formula(form.perf), data = fn)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7905 -0.2756 -0.1563  0.3429  1.0078
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   8.788e-02  3.768e-02   2.332  0.019893
## hits          6.962e-04  1.001e-03   0.695  0.487053
## assists       3.445e-02  1.020e-02   3.377  0.000764
## accuracy     -4.164e-01  1.081e-01  -3.850  0.000126
## head_shots    -4.808e-03  3.149e-03  -1.527  0.127057
## damage_to_players 4.728e-04  5.713e-05   8.275  4.31e-16
##
## (Intercept)      *
```

Comparing models

```
m.games <- lm(as.formula(form.games),fn)
summary(m.games)
```

```
##
## Call:
## lm(formula = as.formula(form.games), data = fn)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.0107  -0.2320  -0.1275   0.2028   0.9583
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.171e-02  2.210e-02   1.888 0.059397
## eliminations   1.151e-02  9.765e-03   1.178 0.238921
## revives        6.993e-02  1.809e-02   3.865 0.000119
## distance_traveled 1.805e-04  1.755e-05  10.287 < 2e-16
## materials_gathered -2.550e-06  3.515e-05  -0.073 0.942186
##
## (Intercept)      .
## eliminations
## revives          ***
```

Comparing models

```
m.context <- lm(as.formula(form.context),fn)
summary(m.context)
```

```
##
## Call:
## lm(formula = as.formula(form.context), data = fn)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.4698 -0.3258 -0.2340  0.5857  0.8553
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.027e+01  2.987e+01   3.022  0.00258
## mental_statesober 1.368e-01  2.933e-02   4.663 3.56e-06
## startTime    -5.672e-08  1.881e-08  -3.015  0.00264
## gameIdSession  1.460e-03  1.463e-03   0.998  0.31863
##
## (Intercept)      **
## mental_statesober ***
## startTime        **
## gameIdSession
```

Evaluate Model Fit

```
cvRes <- NULL
for(i in 1:100) {
  inds <- sample(1:nrow(fn),size = round(nrow(fn)*.8),replace = F)
  train <- fn %>% slice(inds)
  test <- fn %>% slice(-inds)

  # Train
  mTmp.perf <- lm(as.formula(form.perf),train)
  mTmp.games <- lm(as.formula(form.games),train)
  mTmp.context <- lm(as.formula(form.context),train)


  # Test
  toEval <- test %>%
    mutate(prob.p = predict(mTmp.perf,newdata = test),
           prob.g = predict(mTmp.games,newdata = test),
           prob.c = predict(mTmp.context,newdata = test),
           truth = factor(won,levels = c('1','0')))

  auc.p <- roc_auc(toEval,truth,prob.p) %>%
    mutate(model = 'performance')

  auc.g <- roc_auc(toEval,truth,prob.g) %>%
    mutate(model = 'games')
```

Evaluate Model Fit

```
cvRes %>%  
  ggplot(aes(x = .estimate, y = model)) +  
  geom_boxplot() + labs(x = 'AUC', y = 'Specification')
```



Random Forests

```
require(ranger) # Fast random forests package
```

```
## Warning: package 'ranger' was built under R version 4.3.3
```

```
rf.f <- ranger(formula = as.formula(form.full), data = fn)

toEval <- fn %>%
  mutate(prob_won = rf.f$predictions) %>%
  mutate(truth = factor(won, levels = c('1', '0')))

roc_auc(toEval, truth, prob_won)
```

```
## # A tibble: 1 × 3
##   .metric .estimator .estimate
##   <chr>    <chr>         <dbl>
## 1 roc_auc binary         0.838
```


Random Forest Comparison

```
cvRes <- NULL
for(i in 1:100) {
  inds <- sample(1:nrow(fn),size = round(nrow(fn)*.8),replace = F)
  train <- fn %>% slice(inds)
  test <- fn %>% slice(-inds)

  # Train
  mLM.f <- lm(as.formula(form.full),train)
  mRF.f <- ranger(as.formula(form.full),train)

  # Test
  # NEED TO RUN PREDICTION ON RF FIRST
  tmpPred <- predict(mRF.f,test)


  toEval <- test %>%
    mutate(prob.lm = predict(mLM.f,newdata = test),
           prob.rf = tmpPred$predictions,
           truth = factor(won,levels = c('1','0')))

  auc.lm <- roc_auc(toEval,truth,prob.lm) %>%
    mutate(model = 'linear')

  auc.rf <- roc_auc(toEval,truth,prob.rf) %>%
```

Random Forest Comparison

```
cvRes %>%  
  ggplot(aes(x = .estimate, y = model)) +  
  geom_boxplot() + labs(x = 'AUC', y = 'Algorithm')
```



What matters most?

- Random Forests are particularly suitable for investigating **variable importance**
 - I.e., which X predictors are most helpful?
- A few options, but we rely on **permutation tests**
 - Idea: run the best model you have, then re-run it after "permuting" one of the variables
 - "Permute" means randomly reshuffle...breaks relationship
 - How much **worse** is the model when you break a variable?

Variable Importance


- In `ranger()`, use `importance = "permutation"`

```
rf.full <- ranger(formula = as.formula(form.full), data = fn %>%  
  mutate(mental_state = ifelse(mental_state ==  
    'sober', 1, 0)), importance = 'permutation')  
  
rf.full$variable.importance
```

```
##           hits           assists           accuracy  
##    0.013744542    0.003357181    0.004465332  
##    head_shots  damage_to_players  eliminations  
##    0.006363476    0.032384850    0.010595106  
##           revives  distance_traveled  materials_gathered  
##    0.004715412    0.070833478    0.025052451  
##    mental_state      startTime    gameIdSession  
##    0.003392765    0.002135140    0.026763760
```

Variable Importance

```
toplot <- data.frame(vimp = rf.full$variable.importance,  
                     vars = names(rf.full$variable.importance))  
  
toplot %>%  
  ggplot(aes(x = vimp, y = reorder(vars, vimp))) +  
  geom_bar(stat = 'identity') + labs(x = 'VIMP', y = 'Predictor')
```



LASSO

- "Least Absolute Shrinkage and Selection Operator"
- Concept: Make it hard for predictors to matter
 - Practice: λ penalizes how many variables you can include
 - $\sum_{i=1}^n (y_i - \sum_j x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j|$
 - Minimize the errors, but penalize for each additional predictor
 - You *could* kitchen-sink a regression and get super low errors
 - LASSO penalizes you from throwing everything into the kitchen sink
- In **R**, need to install a new package! `install.packages('glmnet')`

```
require(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.3.2
```

LASSO

- Function doesn't use formulas
- Give it the raw data instead, divided into **Y** (outcome) and **X** (predictors)

```
rhsVars <- str_split(gsub('won ~ ', '', form.full), ' \\+ ')[[1]]
fn <- fn %>%
  drop_na(all_of(rhsVars))
Y <- fn$won
X <- fn %>% select(all_of(rhsVars)) %>%
  mutate(mental_state = ifelse(mental_state == 'sober', 1, 0),
         startTime = as.numeric(startTime))
```


LASSO

- Now estimate!

```
lassFit <- glmnet(x = as.matrix(X),  
                  y = as.matrix(Y))
```



LASSO

```
plot(lassFit)
```



Has its own CV!


```
cv.lassFit <- cv.glmnet(x = as.matrix(X), y = as.matrix(Y))  
plot(cv.lassFit)
```



Variable Importance

```
best <- cv.lassFit$glmnet.fit$beta[,cv.lassFit$index[2,]]
vimpLass <- data.frame(vimp = best,
                      vars = names(best))

vimpLass %>%
  ggplot(aes(x = vimp, y = reorder(vars, vimp))) +
  geom_bar(stat = 'identity')
```



Conclusion

- Lots of powerful tools out there!
- Make sure to take **more classes** on these topics!
- Go to Brightspace and take the **16th** quiz
- **Homework:**
 - Problem Set 9 (due 2024-03-29 by 11:59PM)
 - HW 17